

# Exposing memory level parallelism at compile time by **slicing** indirections $C[B[A[i]]]$ and **decoupling** memory accesses.

## DPREF: Decoupling using Dead Code Elimination for Prefetching Irregular Memory Accesses

Konstantinos Sotiropoulos, Jonas Skeppstedt, Angelos Arelakis, Per Stenström  
Chalmers University of Technology, Lund University, ZeroPoint Technologies

### DCE for slicing dead code

Dead Code Elimination (DCE) is a compiler optimization used to remove **ineffec-  
tual** code. “Pre-live” statements like line 6, which affect system state, form the  
slicing basis for DCE.

```
1 void func(int *A, int B) {  
2     int C = B % 5; // data dependency of line 6  
3     if (C == 0)    // control dependency of line 6  
4         B = 0;  
5     else  
6         *A = C;    // "pre-live"  
7 }
```

### Slicing the Sparse Matrix-Vector multiplication

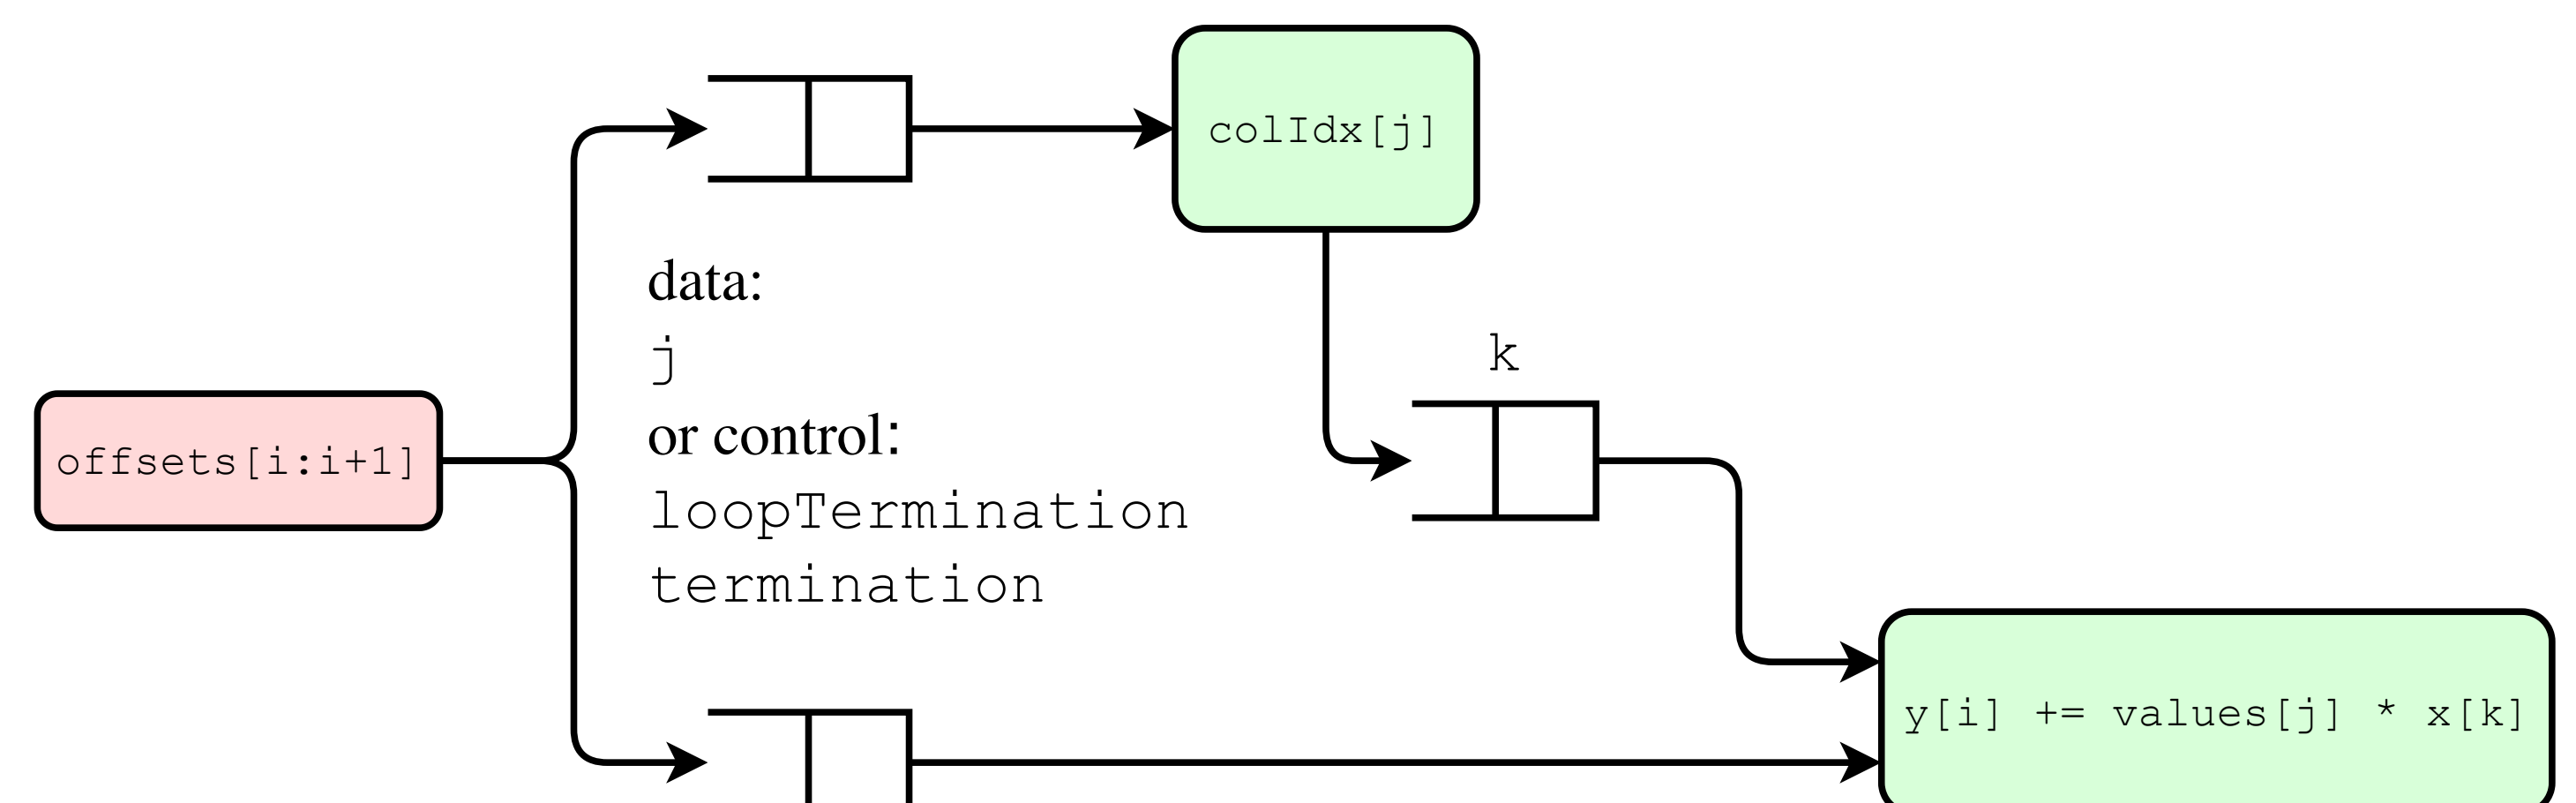
DCE can evolve into a **flexible slicing technique** by redefining what is consid-  
ered as “pre-live”. The `pragma` in line 9 characterizes line 10 as “pre-live”, with  
the dependency on `j` considered satisfied.

```
1 // Compute y = Ax where A is a sparse matrix  
2 // represented by: offsets, colIdx and values  
3 void SpMV(...) {  
4     // for every row of the matrix  
5     for (int i = 0; i < N; i++) {  
6         y[i] = 0;  
7         // for every non-zero column of row i  
8         for (int j = offsets[i]; j < offsets[i+1]; j++)  
9             #pragma dpref sliceOn(j)  
10            y[i] += values[j] * x[colIdx[j]];  
11     }  
12 }
```

Shaded in green is exclusively “live” code, shaded in red is exclusively “dead”  
code, and code in yellow is common to both.

### Slice arrangement

Slices, mapped to threads, are arranged in a dataflow manner using queues for  
inter-slice communication, which carry both data and control tokens.



### Microarchitectural challenges

The slicing requires ISA-visible queues, which should support:

- low-latency, fine-grained and blocking push/pop operations
- low-latency checks for control flow tokens

### Related Work

Manocha, A., Sorensen, T., Tureci, E., Matthews, O., Aragón, J.L., and Martonosi, M. 2021. **GraphAttack: Optimizing Data Supply for Graph Applications on In-Order Multicore Architectures** ACM Transactions on Architecture and Code Optimization (TACO) 18, 1–26.

Nguyen, Q.M. and Sanchez, D. 2023. **Phloem: Automatic Acceleration of Irregular Applications with Fine-Grain Pipeline Parallelism** 2023 IEEE International Symposium on High-Performance Computer Architecture (HPCA).

